

Policy-Grounded RAG for Reliable Enterprise Service Desk Automation: Decision Accuracy, Citation Faithfulness, Safe Abstention, and SLA-Aware Routing

Shruti Sinha

Department of Computer Science and Engineering
Dronacharya College of Engineering, Gurugram, India

Abstract

Enterprise organizations increasingly use large language models (LLMs) to automate internal knowledge access and employee support. Conventional LLM assistants, however, may produce responses that are unsupported by organizational policy, fail to distinguish between departments, or answer confidently when available evidence is insufficient. Retrieval-Augmented Generation (RAG) grounds model output in external knowledge, but recent research shows that retrieval alone does not guarantee faithful attribution, policy compliance, or reliable behavior under conflicting or incomplete evidence. This paper proposes the **Policy-Grounded Enterprise Service Agent (PG-ESA)** for internal service-desk automation. The framework combines tenant-specific policy retrieval, evidence validation, policy citation, LLM-based decision making, confidence-aware safe abstention, department routing, SLA assignment, and audit logging. Rather than producing only a textual reply, the agent selects among operational actions: resolve, route, clarify, escalate, reject, or abstain. A prototype scenario built on a synthetic enterprise dataset (Veridian Corp: ten policies, fifteen employee requests, and historical tickets) illustrates the framework, and an evaluation protocol is defined that compares an ungrounded LLM, a basic RAG system, and the proposed system using decision accuracy, routing accuracy, citation accuracy, unsupported-decision rate, abstention precision and recall, tenant-leakage rate, SLA accuracy, and latency. Quantitative results are intentionally not reported here; they are marked **[TO BE MEASURED]** pending execution of the experiments. The work contributes a practical path from answer-oriented RAG to action-aware, policy-constrained enterprise RAG with human escalation when automated decisions cannot be safely justified.

Keywords: Retrieval-Augmented Generation, Enterprise AI, Policy Grounding, Safe Abstention, Citation Faithfulness, AI Agents, Service Desk Automation, Ticket Routing, SLA, Multi-Tenant Systems.

1. Introduction

Large language models have shown strong capabilities in natural-language understanding, generation, and question answering, which has encouraged organizations to integrate them into internal support systems, enterprise search, and automated service workflows. Enterprise environments, however, differ substantially from open-domain conversation. An internal service agent must answer according to *organization-specific policies*, not according to general knowledge stored in a model's parameters. Whether an employee is eligible for VPN access, whether a software installation requires security approval, or whether an equipment request needs Finance authorization may depend entirely on internal rules.

Retrieval-Augmented Generation (RAG) combines parametric model knowledge with retrieved non-parametric knowledge and improves performance on knowledge-intensive tasks [1]. Yet RAG does not by itself solve enterprise reliability: retrieved evidence can be irrelevant, incomplete, conflicting, or insufficient. Self-RAG addresses part of this by letting the model decide when to retrieve and critique its own output [2]. Enterprise-specific work such as EnterpriseRAG evaluates noisy retrieval, knowledge gaps, factual conflicts, and complex instructions, and reports a gap between satisfying individual constraints and satisfying all requirements at once [10]. COMPASS evaluates organization-specific allowlist and denylist policies and shows that good performance on legitimate requests does not necessarily translate into reliable enforcement of prohibited ones [9].

Attribution adds a further difficulty. A response may cite a document that technically supports its statement without the citation faithfully reflecting how the statement was produced. Wallat et al. distinguish citation

correctness from citation *faithfulness* and report substantial faithfulness problems in attributed RAG output [5]. A reliable enterprise agent therefore needs more than *retrieve, generate, answer*. It needs:

retrieve → *validate evidence* → *reason* → *verify citation* → *determine safe action* → *execute or abstain* → *audit*

This paper proposes such a framework and an evaluation protocol for it. The main contributions are summarized in Section 9.

2. Problem Statement, Research Questions, and Hypotheses

Internal service desks handle requests such as password resets, VPN access, software installation, laptop replacement, printer faults, mailbox quotas, guest Wi-Fi, expense-tool access, phishing reports, administrative access, and home-office equipment. These involve multiple departments (IT, Security, Finance, HR, Asset Management, Management). Conventional ticketing requires staff to interpret policy and route manually. The proposed research addresses eight problems:

Table 1. Problems addressed by the proposed framework.

ID	Problem	Description
P1	Policy fragmentation	Rules are distributed across several policy documents.
P2	Inconsistent interpretation	Different staff may interpret the same policy differently.
P3	Unsupported LLM responses	An LLM can give a plausible answer when no supporting policy exists.
P4	Incorrect routing	A request may go to IT although the governing policy assigns it to Finance or Security.
P5	Insufficient evidence	The request may lack the information needed for a safe decision.
P6	Policy conflict	Overlapping policies may state different conditions (e.g., standard refresh vs. early replacement).
P7	Lack of transparency	Employees may not know which policy justified a decision.
P8	Lack of operational awareness	Text alone does not set ticket state, department, priority, or SLA.

2.1 Research Questions

- **RQ1.** Does policy-grounded RAG improve decision accuracy over an ungrounded LLM for enterprise service-desk requests?
- **RQ2.** Does explicit policy-citation validation reduce unsupported decisions?
- **RQ3.** Can safe abstention identify requests whose available policy evidence is insufficient, ambiguous, or contradictory?
- **RQ4.** Can the framework route requests to the correct department while maintaining tenant isolation?
- **RQ5.** How does operational context (ticket state, SLA information) affect the reliability of automated service decisions?

2.2 Hypotheses

- **H1.** Policy-grounded RAG achieves higher decision accuracy than an ungrounded LLM.
- **H2.** Citation validation reduces unsupported automated decisions.
- **H3.** Safe abstention reduces unsafe automated decisions when evidence is insufficient.
- **H4.** Policy-aware routing improves department-assignment accuracy.
- **H5.** Tenant-level retrieval filtering prevents cross-organization policy leakage.

These hypotheses must be tested experimentally and are not presented as established findings.

3. Related Work

3.1 Retrieval-Augmented Generation

Lewis et al. introduced RAG to combine parametric and retrieved non-parametric knowledge, highlighting benefits for knowledge-intensive tasks as well as provenance and knowledge-update advantages over parametric-only models [1]. It is the foundation of the proposed architecture.

3.2 Enterprise RAG

Enterprise information is proprietary, permission-controlled, heterogeneous, and continuously updated. EK-RAG studies enterprise knowledge question answering over corporate documents [6]; other work addresses Q&A benchmark generation for enterprise RAG [7] and deep search over heterogeneous enterprise data [8]. EnterpriseRAG emphasizes retrieval noise, knowledge gaps, factual conflicts, and complex multi-constraint instructions as deployment failure modes [10]. These findings motivate enterprise-specific evaluation instead of reliance on general QA benchmarks.

3.3 Hallucination in RAG

RAG reduces but does not eliminate unsupported generation. RAGTruth provides roughly 18,000 RAG-generated responses with manual hallucination annotations, showing that unsupported or contradictory claims persist in RAG systems [4]. The proposed system therefore treats retrieval as *evidence acquisition*, not as a guarantee of correctness. Automated RAG evaluation frameworks such as RAGAS [3] offer reference-free metrics that can complement human annotation.

3.4 Citation Faithfulness

Citation-based RAG aims to make responses auditable. Wallat et al. argue that a citation can support a statement without having influenced its generation and report citation-faithfulness problems in a substantial share of citations under their setting [5]. The proposed evaluation therefore checks both whether the cited policy supports the decision and whether the citation corresponds to the evidence actually used by the decision pipeline.

3.5 Self-Reflective RAG and Safe Abstention

Self-RAG uses retrieval and self-reflection to improve factuality and citation accuracy [2]. Abstention matters in high-risk settings because an unsupported answer can be worse than an explicit request for human intervention. Recent work studies evidence sufficiency for answer-abstention calibration in RAG [11], confidence-based response abstinence [12], and self-aware trust-or-abstain RAG [13]. This paper extends the idea from *answer-level* to *action-level* abstention: instead of "I cannot determine whether the request is allowed," the system emits "Automated action withheld; ticket routed to Security/Manager for human review with supporting evidence."

3.6 Research Gap

Table 2. Existing capabilities by research area.

Research area	Existing capability
RAG	Retrieval-grounded generation
Enterprise RAG	Corporate knowledge retrieval and benchmarks
Citation research	Attribution and faithfulness evaluation
Hallucination research	Detection and mitigation
Safe abstention	Refusing or deferring uncertain answers
Policy alignment	Organization-specific allow/deny evaluation

Prior work often evaluates these capabilities independently or in question-answering settings. Work such as EPRAG addresses action selection under multi-source enterprise retrieval and shows that source preference, conflict handling, and abstention are central enterprise problems [14]. The gap addressed here is one of **integration and evaluation**: there is a need for experimentally evaluated service-agent architectures that

connect policy-grounded retrieval with citation validation, action-level abstention, department routing, ticket state, SLA handling, and tenant isolation in one stateful workflow. This is not a claim that each component is individually novel.

4. Proposed Framework: PG-ESA

The Policy-Grounded Enterprise Service Agent processes each request through seven stages, shown in Figure 1: tenant identification, access-controlled policy retrieval, evidence and confidence evaluation, LLM decision, citation validation, an action gate, and SLA assignment with audit logging.




Figure 1. PG-ESA processing pipeline (top row left to right, bottom row right to left).

4.1 Presentation and Roles

Employees submit requests, track status, converse with the agent, and inspect policies. Departments see ticket queues with SLA timers, comments, resolution, and reassignment. Administrators manage policies, users, departments, SLA configuration, and audit logs.

4.2 Multi-Tenant Isolation

Each organization (tenant) owns its users, policies, requests, tickets, and audit logs. Retrieval is constrained by tenant identity:

$$D_q = \{ d \in D \mid \text{tenant}(d) = \text{tenant}(q) \}$$

where D is the full document collection, q the request, and D_q the permitted retrieval set. The LLM must never receive documents outside D_q . The tenant leakage rate is

$$TLR = (N_{\text{cross-tenant}} / N_{\text{retrieval}}) \times 100, \text{ target } TLR = 0\%$$

4.3 Policy Retrieval

Each policy is stored with structured metadata: policy ID, tenant ID, title, responsible department, content, version, effective date, status (active/disabled), and embedding. Policy text is embedded and stored in a vector-enabled PostgreSQL database. For request q with embedding E_q and policy embedding E_d , similarity is

$$\text{sim}(q, d) = (E_q \cdot E_d) / (\|E_q\| \cdot \|E_d\|), R(q) = \text{Top}K_{d \in D_q} \text{sim}(q, d)$$

4.4 Evidence Validation

Retrieval alone is insufficient. The system evaluates relevance, policy status, tenant ownership, applicability, conflicts between policies, and evidence completeness. A conceptual evidence score is

$$C_e = w_1 R + w_2 A + w_3 S + w_4 T$$

where R is retrieval relevance, A applicability, S source/status validity, and T tenant validity. The weights should be determined experimentally on a validation set rather than fixed arbitrarily.

4.5 LLM Decision Engine and Action Space

The LLM receives the request, retrieved policy evidence, relevant ticket context, and the permitted action set, and must return structured output. The system does not permit arbitrary actions.

```
{
  "decision": "ROUTE",
  "department": "IT",
  "policy_ids": ["KB-01"],
  "reason": "Account exceeded five failed attempts.",
  "requires_human": true,
  "confidence": "high"
}
```

Table 3. Action space of the agent.

Action	Name	Used when
A1	Resolve	Policy clearly permits direct resolution.
A2	Route	Another department or human role must handle the request.
A3	Clarify	Required information is missing.
A4	Escalate	Security-sensitive or high-risk cases.
A5	Reject	Request is explicitly prohibited or fails policy requirements.
A6	Abstain	Evidence is insufficient, conflicting, or unauthorized to act on.

4.6 Citation Validation

Every policy-dependent automated decision carries the policy ID, policy title, the relevant evidence, the decision, and the reason. For example: *Decision: manager approval required; Policy: KB-06, Email Mailbox Quota; Evidence: quota increases beyond 25 GB require manager approval and are capped at 50 GB*. Citation correctness is

$$CC = (N_{supported} / N_{citations}) \times 100$$

Citation correctness is not treated as equivalent to citation faithfulness, since a technically correct citation can still be post-rationalized [5]; faithfulness is checked by comparing the cited policy to the evidence actually passed to the decision step.

4.7 Safe Abstention

The system abstains when there is no relevant policy, low evidence confidence, a policy conflict, missing required information, or a high-risk action without authorization. Abstention is operational, not a bare "I don't know":

```
Status: HUMAN_REVIEW_REQUIRED
Reason: Available policy evidence is insufficient
       to authorize the requested action.
Evidence: [retrieved policy information]
Recommended department: [Department]
Required information: [Missing information]
```

4.8 Routing, SLA, and Audit

Routing is modeled as $\text{Department} = f(\text{Request}, \text{Policy}, \text{Evidence}, \text{Risk})$ and evaluated separately from answer quality. SLA assignment, $\text{SLA} = f(\text{Priority}, \text{Department}, \text{Risk}, \text{Policy})$, links the decision layer to workflow execution (for example, a security incident maps to high priority, the Security queue, a short SLA, and escalation on breach). SLA routing is an operational component, not a claimed novelty. Each decision yields an audit record (request ID, tenant, timestamp, retrieved policies, decision, department, citation, confidence, human-review flag, and whether an action was executed), generated independently of the model's prose where possible.

```
{ "request_id": "REQ-08", "tenant_id": "VERIDIAN",
  "retrieved_policies": ["KB-09"], "decision": "ESCALATE",
  "department": "SECURITY", "citation": "KB-09",
  "human_review": true, "action_executed": false }
```

4.9 Guardrail Model

The final action combines evidence sufficiency and authorization. With E denoting evidence sufficiency and τ_r an experimentally calibrated reliability threshold:

Table 4. Guardrail decision rules (threshold calibrated on a validation set).

Condition	Action
$E \geq \tau_r$ AND policy allows	Resolve
$E \geq \tau_r$ AND human approval required	Route
Required information missing	Clarify
High risk OR security incident	Escalate
Policy prohibits the request	Reject
$E < \tau_r$	Abstain (human review)

5. Prototype Dataset

The prototype represents **Veridian Corp** with three sources: a knowledge base, employee requests, and historical tickets.

Table 5. Principal knowledge-base policies (an additional Asset Management policy defines hardware refresh).

ID	Policy	ID	Policy
KB-01	Password Reset	KB-06	Email Mailbox Quota
KB-02	VPN Access	KB-07	Guest Wi-Fi
KB-03	Laptop Replacement	KB-08	Expense Software Access
KB-04	Software Installation	KB-09	Security Incident Reporting
KB-05	Printer Troubleshooting	KB-10	Work-From-Home Equipment

Table 6. Prototype employee requests (15).

ID	Category	Primary issue	Expected handling
REQ-01	Hardware	Dead laptop, 3.5 years old	IT / approval analysis
REQ-02	Wi-Fi	Guest access	Direct / self-service
REQ-03	Password	Six failed attempts	IT
REQ-04	Software	Non-catalog software	Security review
REQ-05	VPN	Expired credentials	IT
REQ-06	Printer	Paper-jam error	IT
REQ-07	Equipment	WFH monitor	Manager + Finance + IT
REQ-08	Security	Phishing email	Security
REQ-09	Mailbox	Full mailbox	Policy-based handling
REQ-10	Access	Finance server admin access	Human approval / review
REQ-11	VPN	Contractor access	Manager approval
REQ-12	Expense	Invalid credentials	Finance / technical support
REQ-13	Hardware	Flickering laptop	Diagnostic assessment
REQ-14	Software	Productivity extension	Policy / security assessment
REQ-15	Ambiguous	"It's not working"	Clarification

The historical ticket queue (TK-1042 to TK-1051) covers resolved, active, closed, rejected, and security-review cases across VPN, laptop, software, mailbox, printer, home-office equipment, phishing, password, admin-access, and guest Wi-Fi issues. Historical tickets serve as *contextual precedent* and are not treated as authoritative policy: **policy is governing evidence; tickets are precedent**.

5.1 Proposed Evaluation Dataset

The 25 available cases suit a prototype demonstration but cannot support broad statistical claims. A stronger evaluation would use roughly 300–500 policy-grounded requests. A proposed 360-case distribution follows; these figures are an *experimental design*, not measured results.

Table 7. Proposed 360-case evaluation distribution.

Category	Cases	Category	Cases
Password	30	Security	35
VPN	30	WFH equipment	30
Laptop	40	Unknown / out-of-policy	30
Software	40	Conflicting / ambiguous	20
Printer	30	Guest Wi-Fi	20
Mailbox	30	Expense	25
		Total	360

Each case should carry a request ID, request text, expected action, expected department, expected policy and citation, human-escalation flag, and SLA class. Ground truth must be created or verified by a human, not accepted automatically from an LLM.

6. Experimental Design

6.1 Systems Compared

- **Baseline 1 – Ungrounded LLM:** Request → LLM → Decision, with no policy retrieval.

- **Baseline 2 – Basic RAG:** Request → Vector retrieval → LLM → Decision, with no citation validator or action gate.
- **Proposed PG-ESA:** Tenant filter → policy retrieval → evidence validation → LLM decision → citation validation → confidence guardrail → action gate → SLA and audit.

6.2 Metrics

Table 8. Evaluation metrics.

Metric	Definition
Decision Accuracy (DA)	Correct decisions / total decisions × 100
Routing Accuracy (RA)	Correct department assignments / routable requests × 100
Citation Accuracy (CA)	Supported citations / total citations × 100
Unsupported Decision Rate (UDR)	Unsupported decisions / automated decisions × 100 (lower is better)
Safe-Abstention Precision (SAP)	Correct abstentions / all abstentions × 100
Safe-Abstention Recall (SAR)	Unsafe cases correctly identified / all unsafe cases × 100
Tenant Leakage Rate (TLR)	Cross-tenant retrievals / total retrievals × 100 (target 0)
SLA Accuracy (SA)	Correct SLA assignments / SLA cases × 100
Latency	End-to-end response time in ms (mean, median, SD, IQR)

6.3 Ablation Study

Table 9. Component ablation configurations.

Config.	RAG	Citation	Guardrail	Abstention	Tenant filter
A	No	No	No	No	No
B	Yes	No	No	No	Yes
C	Yes	Yes	No	No	Yes
D (full)	Yes	Yes	Yes	Yes	Yes

Further single-component removals from configuration D (without citation validation, without confidence guardrails, without abstention, without tenant filtering) quantify each mechanism's contribution.

6.4 Statistical Analysis

Results should not be reported as bare percentages. For paired request-level binary correctness outcomes on the same test cases, McNemar's test is appropriate. Latency is summarized by mean, median, standard deviation, and interquartile range. Confidence intervals should be reported, and statistical significance should not be claimed without a supporting test.

7. Illustrative Case Analysis

7.1 Password Lockout (REQ-03)

Request: "I'm locked out of my account, I tried my password six times." KB-01 states that employees can reset their own password and that accounts locked after five failed attempts require IT assistance. **Expected action: Route to IT.** This tests straightforward policy grounding.

7.2 Phishing (REQ-08)

Request: "I think I got a phishing email." KB-09 governs security incident reporting. **Expected action: Escalate to Security.** This is a high-risk case where ordinary forwarding is inappropriate and policy-based security routing takes priority.

7.3 Guest Wi-Fi (REQ-02)

Request: "Can I get Wi-Fi access for a guest?" KB-07 applies. **Expected handling: direct or self-service resolution;** the system should not create an unnecessary IT ticket.

7.4 Policy Conflict: Laptop Replacement (REQ-01)

The employee reports a completely dead laptop about 3.5 years old. KB-03 indicates eligibility after three years, or earlier with verified hardware failure. The Asset Management policy describes a four-year refresh cycle and requires Finance sign-off, in addition to IT approval, for early replacement outside the cycle. The system must recognize both pieces of evidence (Evidence A: 3-year eligibility / hardware failure; Evidence B: 4-year refresh / Finance sign-off) and identify their relationship instead of selecting whichever passage has the highest embedding similarity. This illustrates *policy-aware conflict handling* and why retrieval alone can be insufficient.

7.5 Out-of-Policy Request

Request: "Can the company pay for my gym membership?" No supplied policy addresses this. A conventional LLM may invent a plausible answer. The proposed system should output: **Decision: ABSTAIN. Reason: no applicable organizational policy retrieved. Action: Human/HR review required. Policy citation: none.** This is a key safe-abstention test.

8. Results

No experiments have yet been run for this draft, so no quantitative results are reported. The table below defines the reporting format; every cell is **[TO BE MEASURED]** and must be filled only from actual experimental output.

Table 10. Results template. TBM = TO BE MEASURED.

Metric	Ungrounded LLM	Basic RAG	Proposed PG-ESA
Decision Accuracy	[TBM]	[TBM]	[TBM]
Routing Accuracy	[TBM]	[TBM]	[TBM]
Citation Accuracy	N/A	[TBM]	[TBM]
Unsupported Decision Rate	[TBM]	[TBM]	[TBM]
Abstention Precision / Recall	N/A	[TBM]	[TBM]
Tenant Leakage Rate	N/A	[TBM]	[TBM]
SLA Accuracy	[TBM]	[TBM]	[TBM]
Average Latency (ms)	[TBM]	[TBM]	[TBM]

For the run, one per-case log should be kept with these fields: *id*, *system*, *predicted_action*, *predicted_department*, *cited_policy*, *evidence_score*, *correct*, *citation_correct*, *abstained*, *latency_ms*. That log supports all metrics, statistical tests, and graphs.

9. Error Analysis, Security, and Discussion

9.1 Error Taxonomy

Table 11. Error categories for post-hoc analysis.

Error type	Description
Retrieval error	Relevant policy not retrieved
Ranking error	Correct policy retrieved but ranked too low
Interpretation error	LLM misinterprets the policy
Citation error	Incorrect policy cited
Hallucination	Unsupported information generated
Routing error	Wrong department assigned
Abstention error	System answers when it should abstain
Over-abstention	System abstains when evidence is sufficient
Tenant leakage	Cross-tenant information retrieved
Policy conflict	Conflicting policies handled incorrectly

9.2 Security Considerations

- **Cross-tenant leakage:** a Tenant A user must never receive Tenant B's policies.
- **Prompt injection:** content in user messages or documents must not override access restrictions.
- **Retrieval poisoning:** an attacker may try to insert misleading policy content.
- **Unauthorized actions:** the LLM must not perform sensitive operations without deterministic authorization checks.
- **Audit manipulation:** audit records should be generated independently of the model's natural-language output.

9.3 Layered View of Correctness

Enterprise reliability should not be measured only by whether the answer reads well. The agent must be able to safely justify and operationalize a decision, which involves seven levels: retrieval, evidence applicability, reasoning, citation, action, authorization, and workflow correctness. This layered view matches enterprise RAG findings on noisy retrieval, knowledge gaps, conflicting evidence, and complex instructions [10].

9.4 Contributions

- **Policy-grounded enterprise agent:** links organization-specific policy retrieval to operational service-desk decisions.
- **Evidence-based decision layer:** ties each automated decision to retrieved policy evidence.
- **Action-level abstention:** extends abstention from text answers to operational actions.
- **Policy-aware routing:** maps policy interpretation to departmental ticket routing.
- **Multi-tenant isolation:** tenant-specific retrieval constraints prevent policy leakage.
- **SLA-aware workflow:** connects decisions to service-level management.
- **Evaluation framework:** compares ungrounded LLM, basic RAG, and policy-grounded RAG with guardrails.

10. Limitations and Future Work

10.1 Limitations

- **Synthetic data:** the Veridian dataset cannot capture the complexity of a real organization.
- **Small prototype set:** 15 requests and the ticket history are insufficient for statistical generalization.
- **Model dependence:** results may vary with the selected LLM.

- **Retrieval dependence:** poor retrieval causes downstream reasoning failures.
- **Policy ambiguity:** real organizations have undocumented exceptions and informal practice.
- **Human evaluation:** citation faithfulness and action appropriateness may need expert judgment.
- **Security evaluation:** production use would need penetration testing and adversarial evaluation.

10.2 Future Work

- Policy versioning and effective-date reasoning; explicit policy-authority and precedence mechanisms for conflicts.
- Learning from human resolutions to improve retrieval and routing; adaptive retrieval depth by request complexity.
- Multimodal evidence (screenshots, documents) and integration with ITSM platforms such as Jira or ServiceNow.
- An open enterprise service-desk benchmark with answerable, unanswerable, conflicting, and adversarial requests.
- Formal, deterministic action verification before any sensitive operation.

11. Conclusion

This paper presented PG-ESA, a policy-grounded framework for improving the reliability of LLM-based internal service-desk automation. It moves RAG from answer generation toward evidence-driven operational decisions: retrieving tenant-specific policies, validating evidence, generating cited decisions, choosing among resolve, route, clarify, escalate, reject, and abstain, assigning SLA information, and keeping an audit trail. An enterprise agent should not merely produce a plausible reply; it should determine whether a response is supported, whether an action is authorized, and whether human intervention is required. The proposed protocol compares an ungrounded LLM, basic RAG, and the full architecture with ablations. Quantitative conclusions should be drawn only after these experiments are executed. The broader direction is **from RAG-based answers to policy-grounded, evidence-supported, safely executable enterprise decisions**.

References

- [1] P. Lewis et al., "Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks," *Advances in Neural Information Processing Systems (NeurIPS)*, 2020.
- [2] A. Asai, Z. Wu, Y. Wang, A. Sil, and H. Hajishirzi, "Self-RAG: Learning to Retrieve, Generate, and Critique through Self-Reflection," *International Conference on Learning Representations (ICLR)*, 2024.
- [3] S. Es, J. James, L. Espinosa-Anke, and S. Schockaert, "RAGAS: Automated Evaluation of Retrieval Augmented Generation," arXiv:2309.15217, 2023.
- [4] C. Niu et al., "RAGTruth: A Hallucination Corpus for Developing Trustworthy Retrieval-Augmented Language Models," in *Proc. ACL*, 2024.
- [5] J. Wallat, M. Heuss, M. de Rijke, and A. Anand, "Correctness is not Faithfulness in Retrieval Augmented Generation Attributions," in *Proc. ICTIR*, 2025, pp. 22–32. DOI: 10.1145/3731120.3744592.
- [6] T. Yu et al., "EKRAG: Benchmark RAG for Enterprise Knowledge Question Answering," in *Proc. KnowledgeNLP*, 2025, pp. 152–159.
- [7] S. Filice et al., "Generating Q&A Benchmarks for RAG Evaluation in Enterprise Settings," in *Proc. ACL Industry Track*, 2025, pp. 469–484.
- [8] P. K. Choubey et al., "Benchmarking Deep Search over Heterogeneous Enterprise Data," in *Proc. EMNLP Industry Track*, 2025, pp. 501–517.
- [9] D. Choi et al., "COMPASS: A Framework for Evaluating Organization-Specific Policy Alignment in LLMs," in *Proc. ACL*, 2026, pp. 46087–46133. DOI: 10.18653/v1/2026.acl-long.2139.
- [10] H. Miao et al., "EnterpriseRAG: Benchmarking LLM Instruction Adherence and Robustness under Non-Ideal Enterprise Retrieval," arXiv, 2026.

- [11] H. Zhang and W. Wu, "Do LLMs Know When Evidence is Insufficient? An Evidence Sufficiency Benchmark for Answer-Abstention Calibration in Retrieval-Augmented Generation," *Computers, Materials & Continua*, vol. 89, no. 1, 2026. DOI: 10.32604/cmc.2026.086343.
- [12] Z. Huang et al., "Confidence-Based Response Abstinence: Improving LLM Trustworthiness via Activation-Based Uncertainty Estimation," in *Proc. UncertainNLP*, 2025.
- [13] X. Zhu et al., "Trust or Abstain? A Self-Aware RAG Approach," arXiv, 2026.
- [14] "EPRAG: An epistemic policy framework for action selection in multi-source enterprise retrieval-augmented generation," *Knowledge-Based Systems*, 2026 (authors to be confirmed).
- [15] "From Documents to Decisions: Enterprise-Grade LLM Systems for Zero-Hallucination, Attributed Generation, and Regulatory Alignment," 2026 (review; authors and venue to be confirmed).